

# Solution Report

## 1 Implementation Details

I implemented a **Sequential Hash Join** (Left-Deep) to evaluate the  $k$ -line join query. This approach builds the result incrementally by joining the accumulated result with the next relation in the sequence. The code can be found in `Solution.js`.

### 1.1 Algorithm Steps

The algorithm proceeds in steps to process a query involving  $k$  relations:

$$R_1 \bowtie R_2 \bowtie \dots \bowtie R_k$$

1. **Step 1:** Join  $R_1$  and  $R_2$  on attribute  $A_2$ .

$$R_{1-2} = R_1 \bowtie_{A_2} R_2$$

I used the Hash Join implementation from Problem 1: Hash  $R_2$  on  $A_2$ , then probe with  $R_1$ .

2. **Step 2:** Join the intermediate result  $R_{1-2}$  with  $R_3$  on attribute  $A_3$ .

$$R_{1-3} = R_{1-2} \bowtie_{A_3} R_3$$

Hash  $R_3$  on  $A_3$ , then probe with  $R_{1-2}$ .

3. **Step  $i$ :** Join the intermediate result  $R_{1-i}$  with  $R_{i+1}$  on attribute  $A_{i+1}$ .

$$R_{1-(i+1)} = R_{1-i} \bowtie_{A_{i+1}} R_{i+1}$$

This “left-deep” strategy means we always join the current intermediate result (which can be large) with a base relation.

## 2 Comparison with Yannakakis Algorithm

Unlike the Yannakakis algorithm (Problem 2), this approach **does not** guarantee global consistency before joining.

- **Intermediate Results:** We might generate intermediate tuples that are later discarded because they don’t match with subsequent relations.
- **Complexity:** In the worst case, the size of intermediate results can be much larger than the input or output, potentially leading to exponential complexity  $O(N^k)$  if the join graph is cyclic or if selectivity is poor. In contrast, Yannakakis guarantees  $O(N + OUT)$ .

## 3 Execution Results

I ran the sequential algorithm with the same datasets used in Problem 2.

**Case 1:**  $k = 3$

- **Input:** 3 relations (same dataset as Problem 2).
- **Result:** Found **4 tuples**.
- **Output:** Matches the output from the Yannakakis implementation.

**Case 2:**  $k = 5$

- **Input:** 5 relations (chain with noise).
- **Result:** Found **15 tuples**.
- **Output:** Matches the output from the Yannakakis implementation.

The sequential hash join correctly computes the result, but it constructs full intermediate tuples at each step.